PRATHEEK B

Gen AI Engineer · Agentic AI & LLM Systems · Product Management Enthusiast
Bengaluru, India · bpratheek633@gmail.com · linkedin.com/in/bpratheek

SUMMARY

Gen AI engineer at Bosch, building production agentic systems now running across 10 OEM programs. I look for the friction before I look for the tool: inconsistent AI outputs, slow documentation search, messy handoffs, then ship fixes teams actually use. Outside work, I'm building an out-of-home advertising marketplace on my own, from customer interviews through the PRD to the full stack. Claude Code and Cursor are part of my everyday workflow, and I hold an IREB CPRE certification in requirements and product thinking.

EXPERIENCE

Gen AI Engineer · Bosch Global Software Technologies, Bengaluru *Jan 2024 – Present*

- Engineers across teams were getting inconsistent Copilot completions because agents had no shared context on team-specific conventions. Built and shipped a VS Code extension (.vsix) that auto-fetches domain-specific instruction sets, now used by 250+ developers, improving prompt consistency by 90% and agent output delivery accuracy by 50%.
- Design, code, and test handoffs were slow and inconsistent across teams. Built a multi-agent SDLC assist system (Code Review Assist, Design Assist, Unit Test Assist) with human-in-the-loop checkpoints, cutting manual review effort by 35% for the teams that adopted it. Built with LangChain, LangGraph, and Azure OpenAI.
- Long multi-agent chains kept losing context between steps, causing agents to act on stale or incorrect information. Fixed this with a layered memory architecture, an immutable intent layer, an append-only decisions log, and delta tracking with reasoning provenance, cutting context-drift errors by 40%.
- Manual test design was slow and inconsistent across reviewers. Built a Knowledge Graph-backed LLM framework for automated test design generation, with memory and feedback loops that improved test design accuracy by 30% across iterations.
- Traced recurring complaints about poor LLM output quality back to inconsistent prompts across teams. Built and shipped a VS Code Prompt Library plugin, adopted by 50+ engineers across the org, cutting prompt iteration time by 45%.
- Engineers were losing significant time manually searching scattered embedded software documentation, causing repeated context switching. Designed and deployed a production RAG pipeline (chunking, retrieval, reranking) on Azure OpenAI and vector databases, integrated via REST APIs, cutting manual search time by 40% and handling 500+ queries a week.
- These tools are running in production across 10 OEM programs, contributing to a 15% efficiency gain and a 25% increase in AI tool adoption across Q1 and Q2 2026.

Project Trainee, Embedded Systems & Test Automation · Bosch Global Software Technologies, Bengaluru *Oct 2022 – May 2023*

- Automated 300+ release test cases in Python and Perl. Embedded systems knowledge from this role fed directly into later AI and RAG pipeline work.

PROJECTS & VENTURES

OOH Marketplace · Independent Project, Early Stage *2024 – Present*

- Advertisers and agencies had no straightforward way to discover and book out-of-home ad space. Ran roughly 20 structured discovery interviews to validate the pain point before writing any code.
- Wrote the full PRD independently, covering 11 sections, 5 personas, 66 functional requirements, and C4 architecture, then built the product full stack (FastAPI backend, React/Vite frontend), moving between the two depending on what the problem needed.
- Building this alongside my full-time role, currently pre-launch and iterating toward an MVP.

LLM Test Mutation Observability Dashboard · Bosch Global Software Technologies, Bengaluru *2026 – Present*

- Test engineers had no visibility into whether LLM-generated mutants were improving real test quality or just adding noise. Built an observability dashboard on Grafana and Prometheus, backed by OpenTelemetry, to make that visible.
- Surfaces token consumption, latency, mutation bias, code coverage, mutation quality, an LLM trust score, coverage gaps, and LLM-as-judge generated recommendations.
- Instrumented the full pipeline: a code check-in triggers GitHub Actions, which compiles the code, generates mutants via the Claude API, runs G-Test against them, and logs results back into the dashboard.

Agentic SDLC Orchestrator · Bosch Internal *2024 – Present*

- Built to close the gap between design docs and shipped code: an end-to-end multi-agent pipeline from design ingestion through code generation, unit testing, and a human review gate. Built with LangChain, Azure OpenAI, and custom orchestration logic.

SKILLS

AI-Native Building: Claude Code, Cursor, agent orchestration, MCPs, skills-based workflows, LangChain, LangGraph, Copilot Agents, prompt engineering at scale

Product: Customer discovery, PRD authoring (IREB CPRE), prototyping & design, shipping independent projects, driving adoption

Engineering: Python, FastAPI, React/Vite, REST APIs, RAG, Knowledge Graphs, Vector DBs, Azure OpenAI, AWS Bedrock (familiar), Docker, CI/CD, Git, Grafana, Prometheus, OpenTelemetry, GitHub Actions

CERTIFICATIONS & EDUCATION

IREB CPRE FL (Apr 2026) · Azure AI Fundamentals AZ-900 · ISTQB-CTFL · ISTQB-AuT

B.Tech, Electrical and Electronics Engineering, REVA University, Bengaluru (2019–2023), CGPA 8.6/10